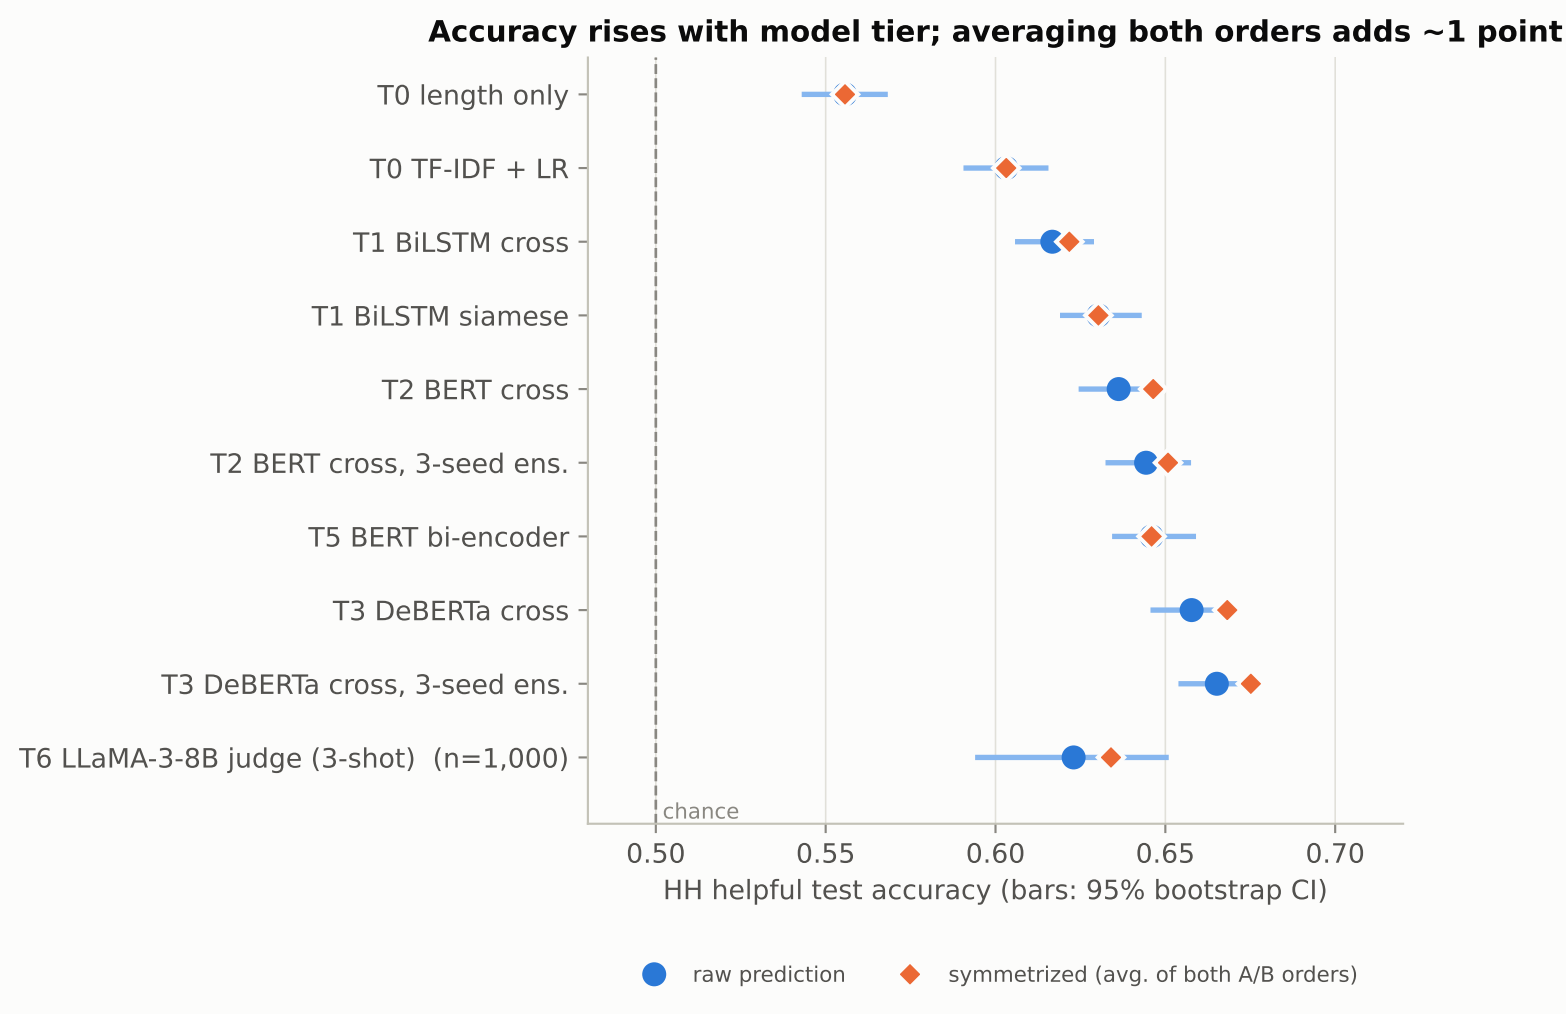
Accuracy rises with model tier; averaging both orders adds ~1 point
T0 length only
T0 TF-IDF + LR
T1 BiLSTM cross
T1 BiLSTM siamese
T2 BERT cross
T2 BERT cross, 3-seed ens.
T5 BERT bi-encoder
T3 DeBERTa cross
T3 DeBERTa cross, 3-seed ens.
T6 LLaMA-3-8B judge (3-shot) (n=1,000)
chance
0.50
0.55
0.60
0.65
0.70
HH helpful test accuracy (bars: 95% bootstrap CI)
raw prediction
symmetrized (avg. of both A/B orders)