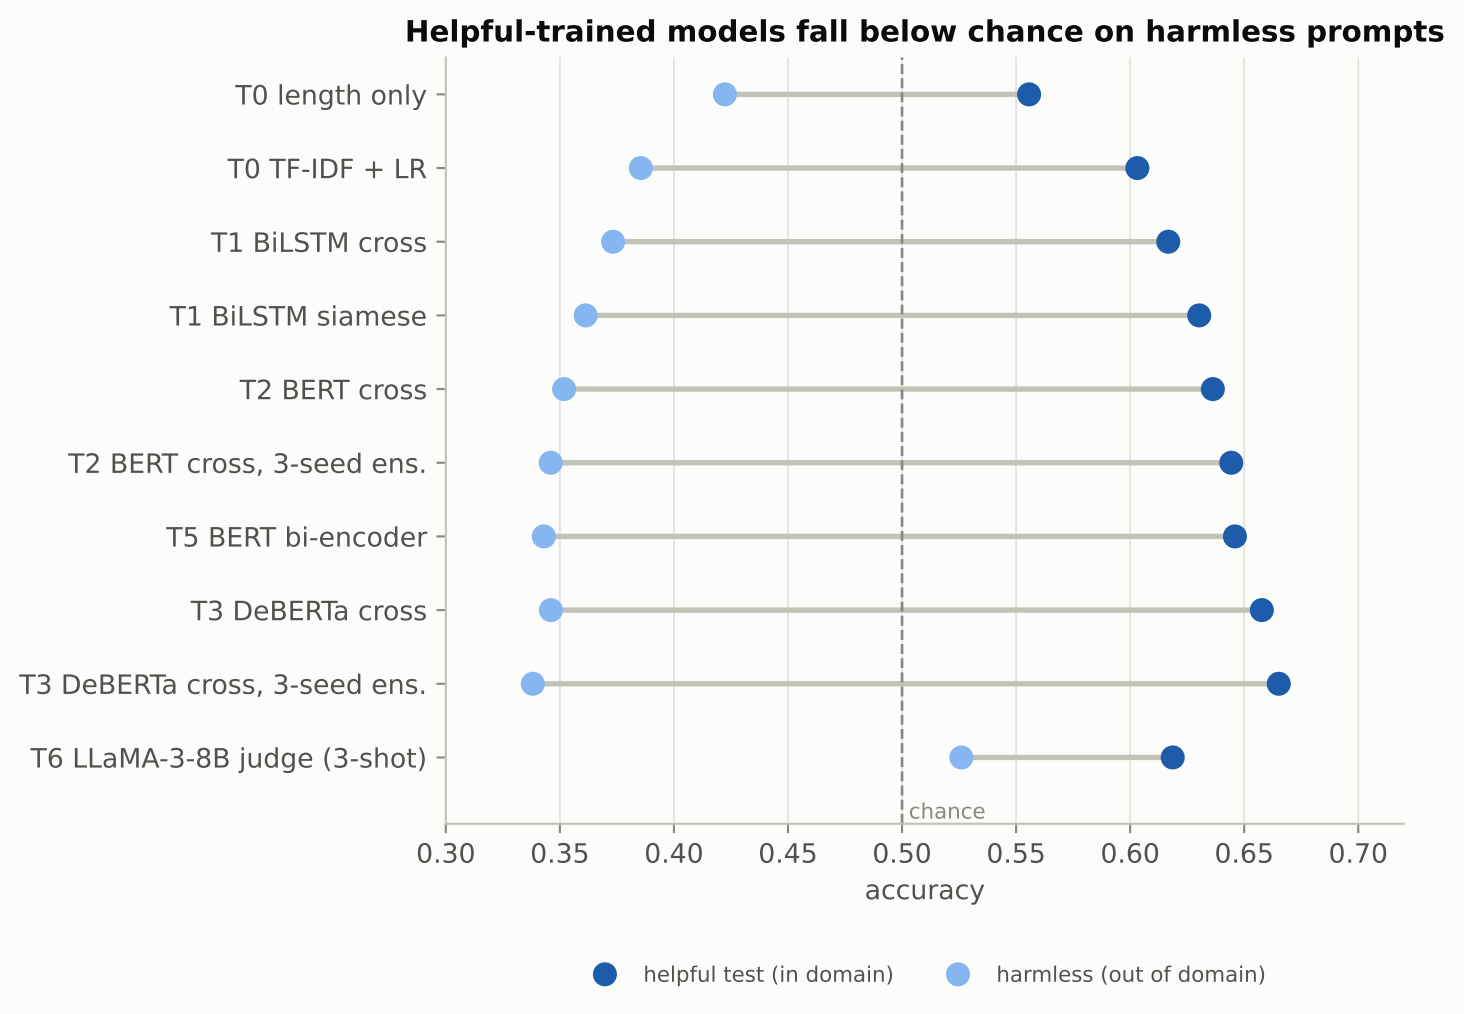
Helpful-trained models fall below chance on harmless prompts
T0 length only
T0 TF-IDF + LR
T1 BiLSTM cross
T1 BiLSTM siamese
T2 BERT cross
T2 BERT cross, 3-seed ens.
T5 BERT bi-encoder
T3 DeBERTa cross
T3 DeBERTa cross, 3-seed ens.
T6 LLaMA-3-8B judge (3-shot)
chance
0.30
0.35
0.40
0.45
0.50
0.55
0.60
0.65
0.70
accuracy
helpful test (in domain)
harmless (out of domain)